

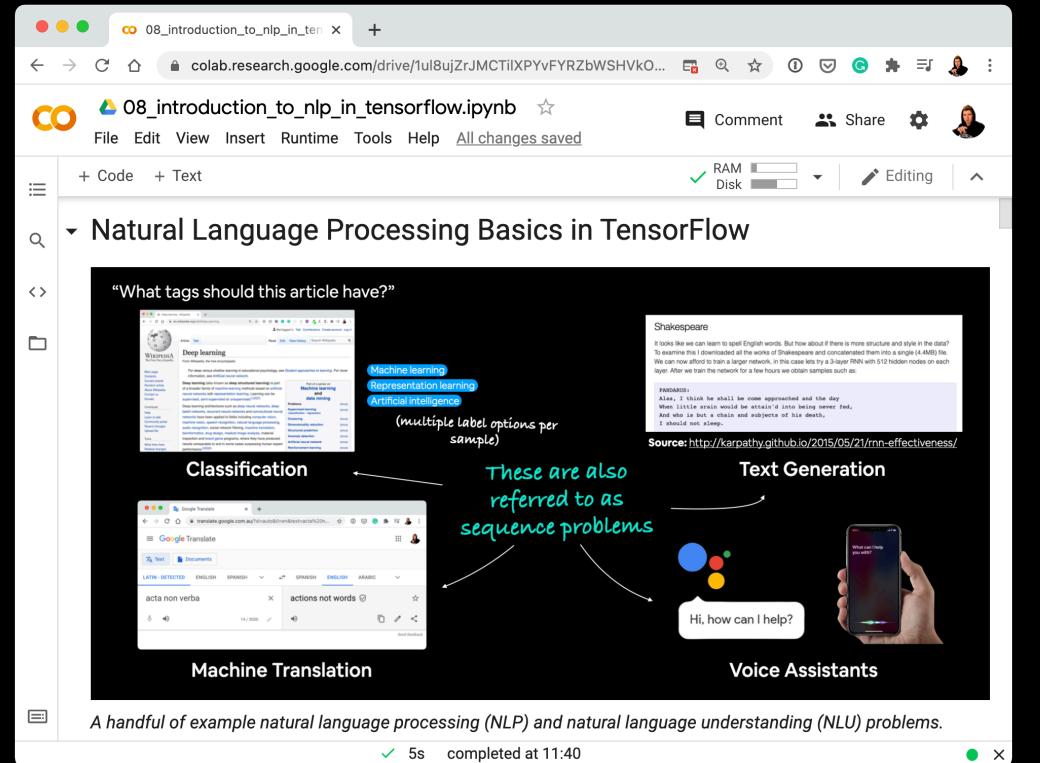
Natural Language Processing (NLP) with



TensorFlow

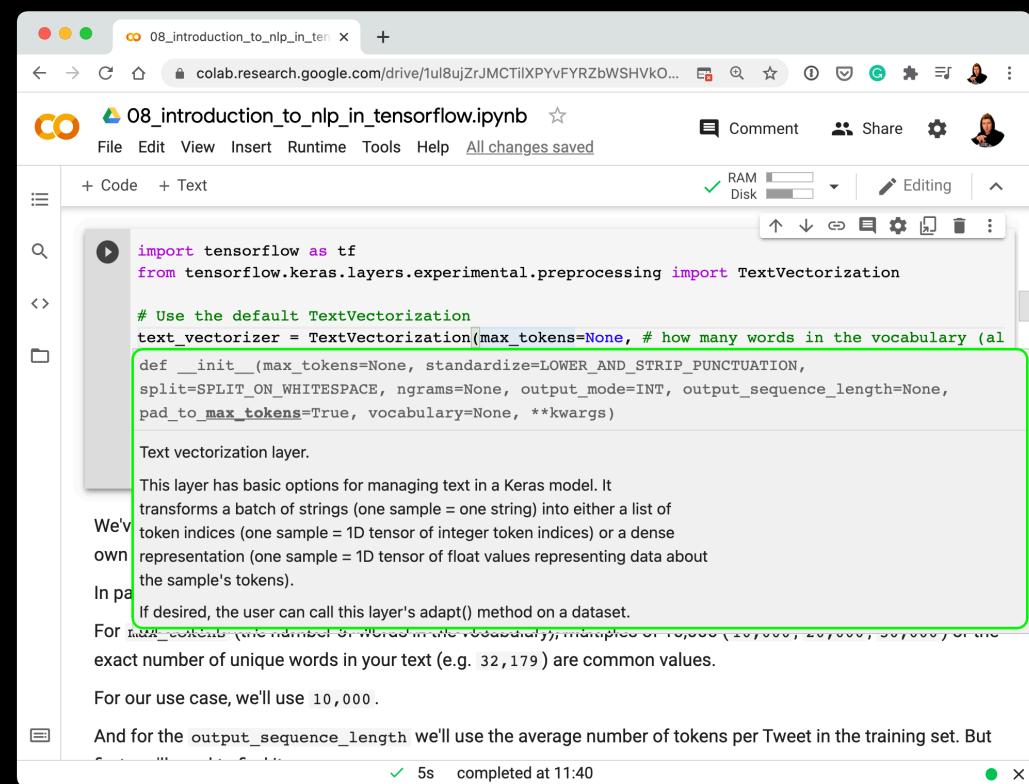
Where can you get help?

- Follow along with the code

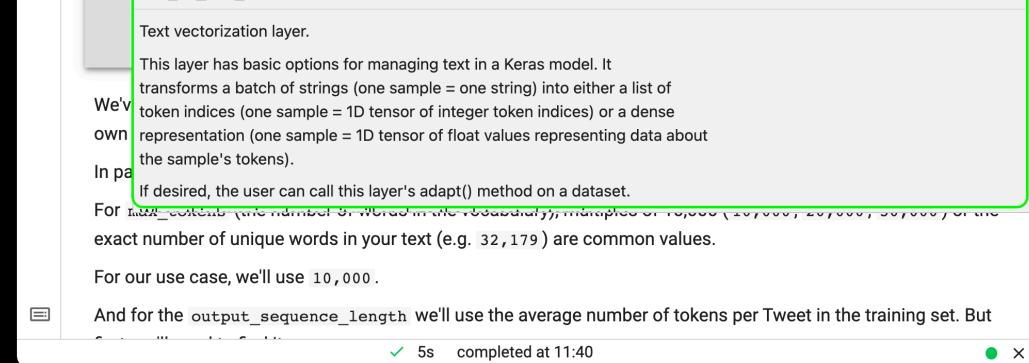


"If in doubt, run the code"

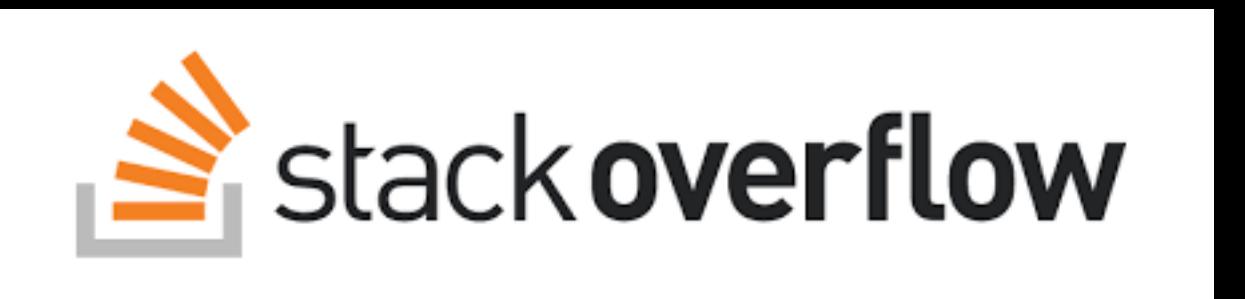
- Try it for yourself



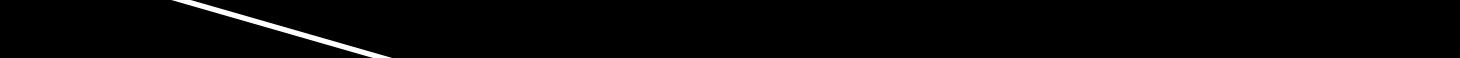
- Press SHIFT + CMD + SPACE to read the docstring



- Search for it

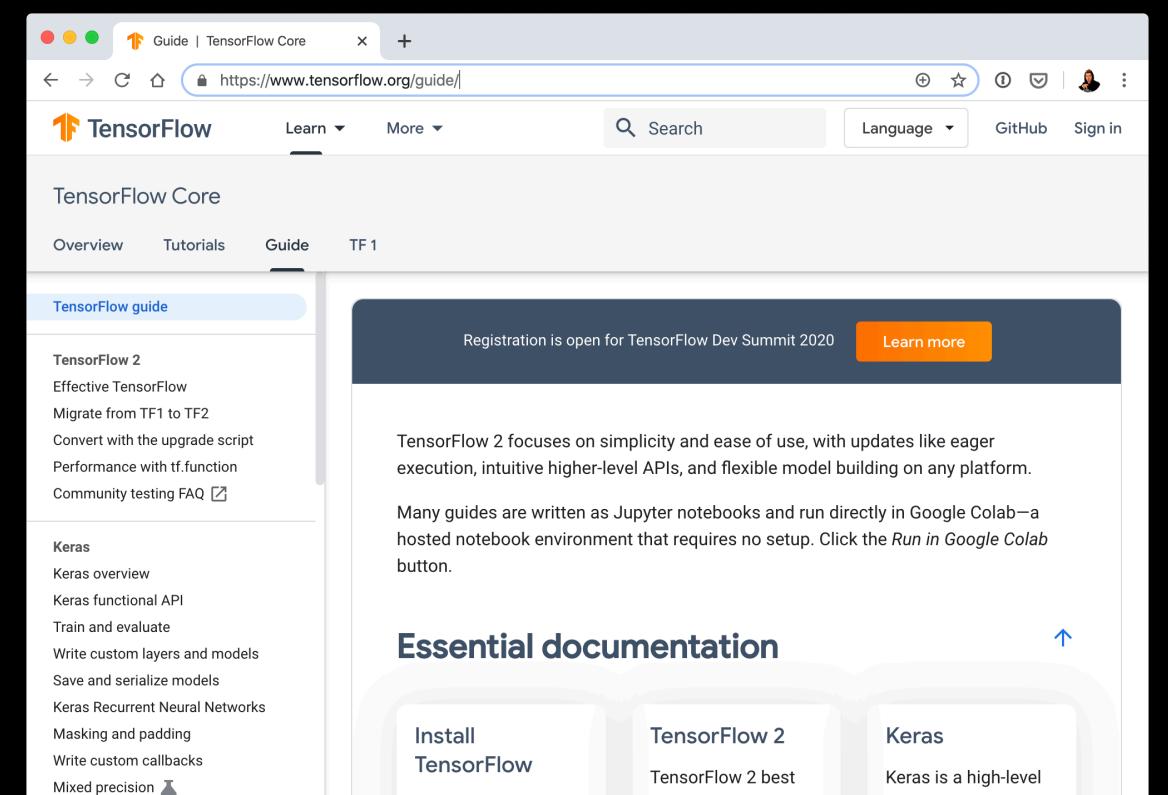


- Try again



- Ask (don't forget the Discord chat!)

(yes, including the "dumb" questions)



“What is a NLP problem?”

Example NLP problems

and NLU... (natural language understanding)

“What tags should this article have?”

The screenshot shows the Wikipedia article on "Deep learning". On the right side of the page, there is a sidebar with a section titled "Machine learning" which includes "Representation learning" and "Artificial intelligence". Below this, a note states "(multiple label options per sample)".

Classification

The screenshot shows the Google Translate interface translating the Latin phrase "acta non verba" into English as "actions not words".

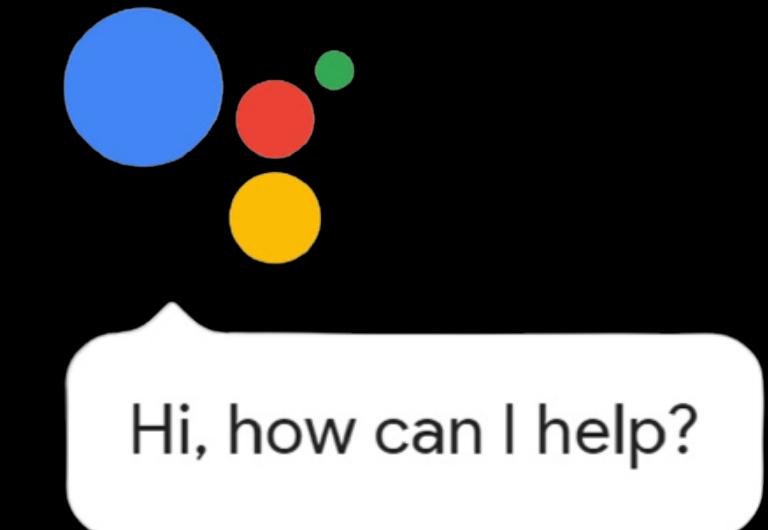
Machine Translation

These are also referred to as sequence problems

The screenshot shows a text generation example from William Shakespeare's "Pandarus". The text reads: "Alas, I think he shall be come approached and the day When little strain would be attain'd into being never fed, And who is but a chain and subjects of his death, I should not sleep."

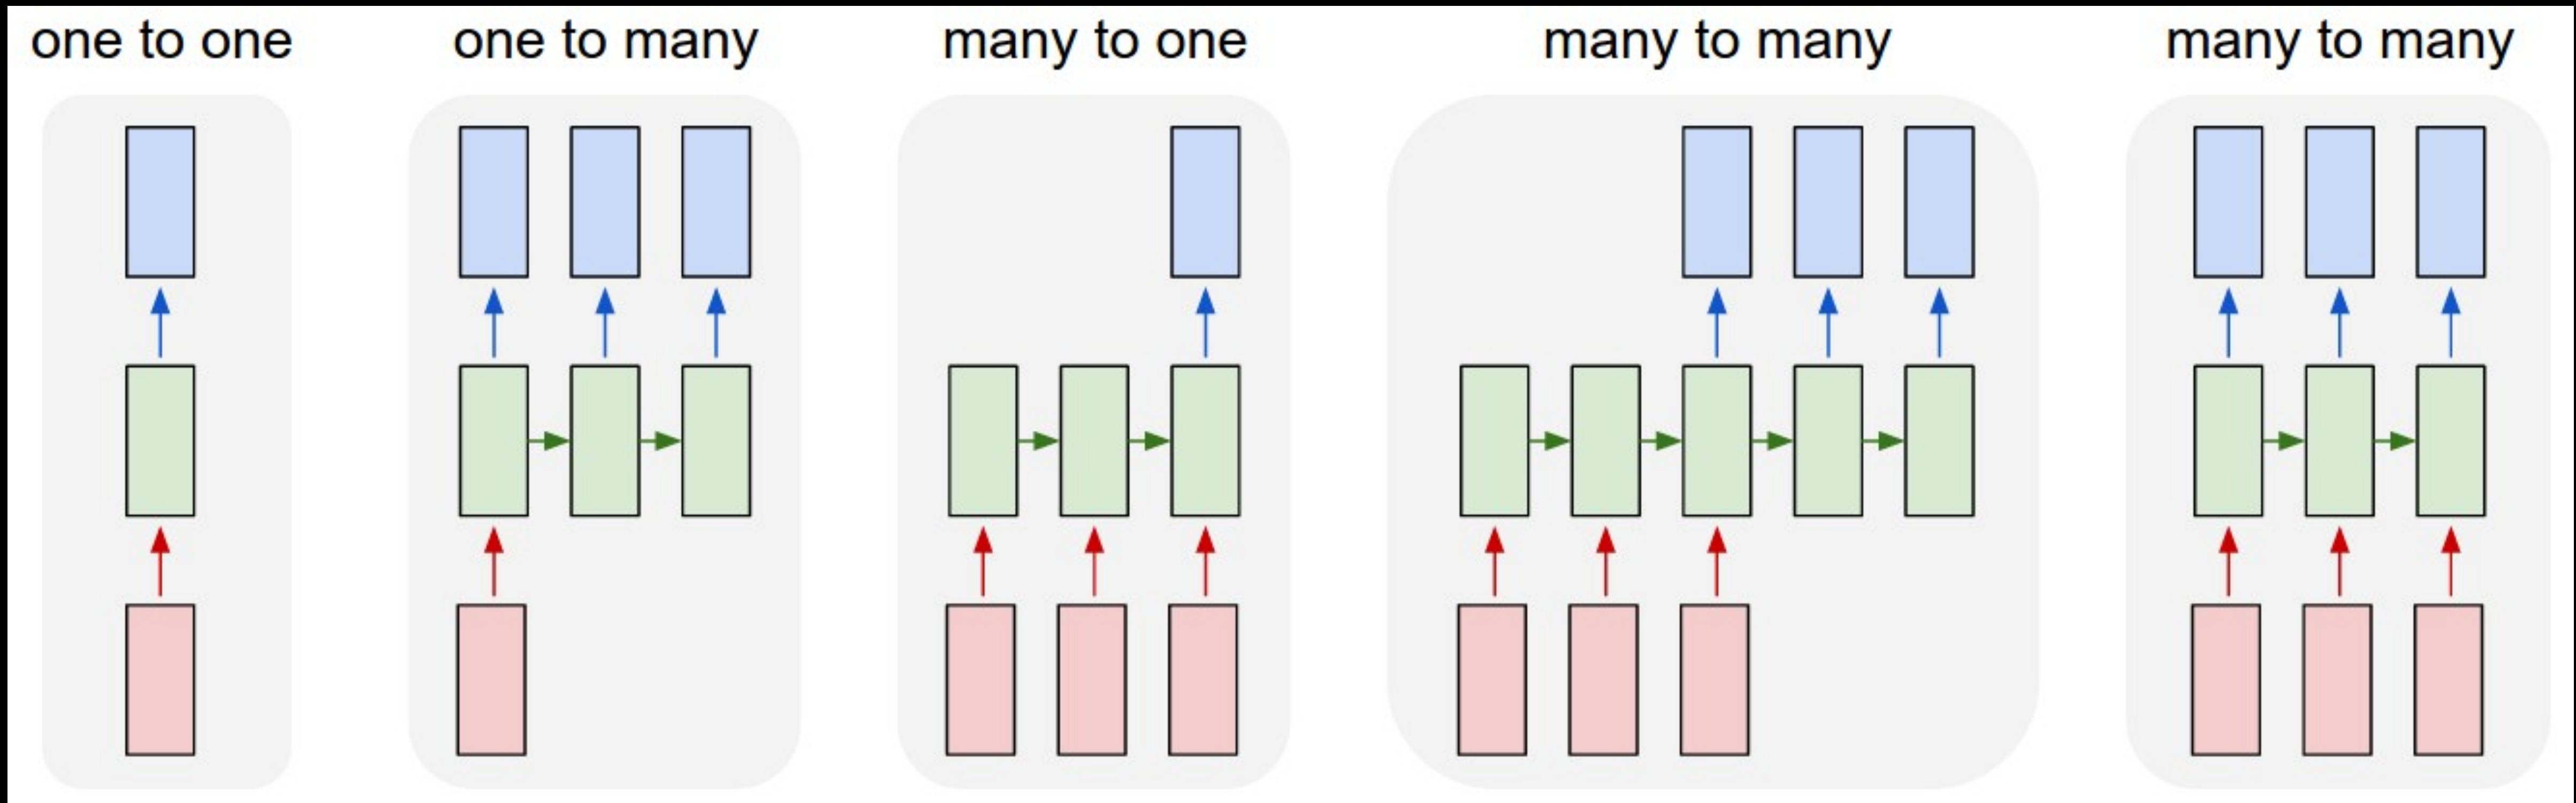
Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Text Generation



Voice Assistants

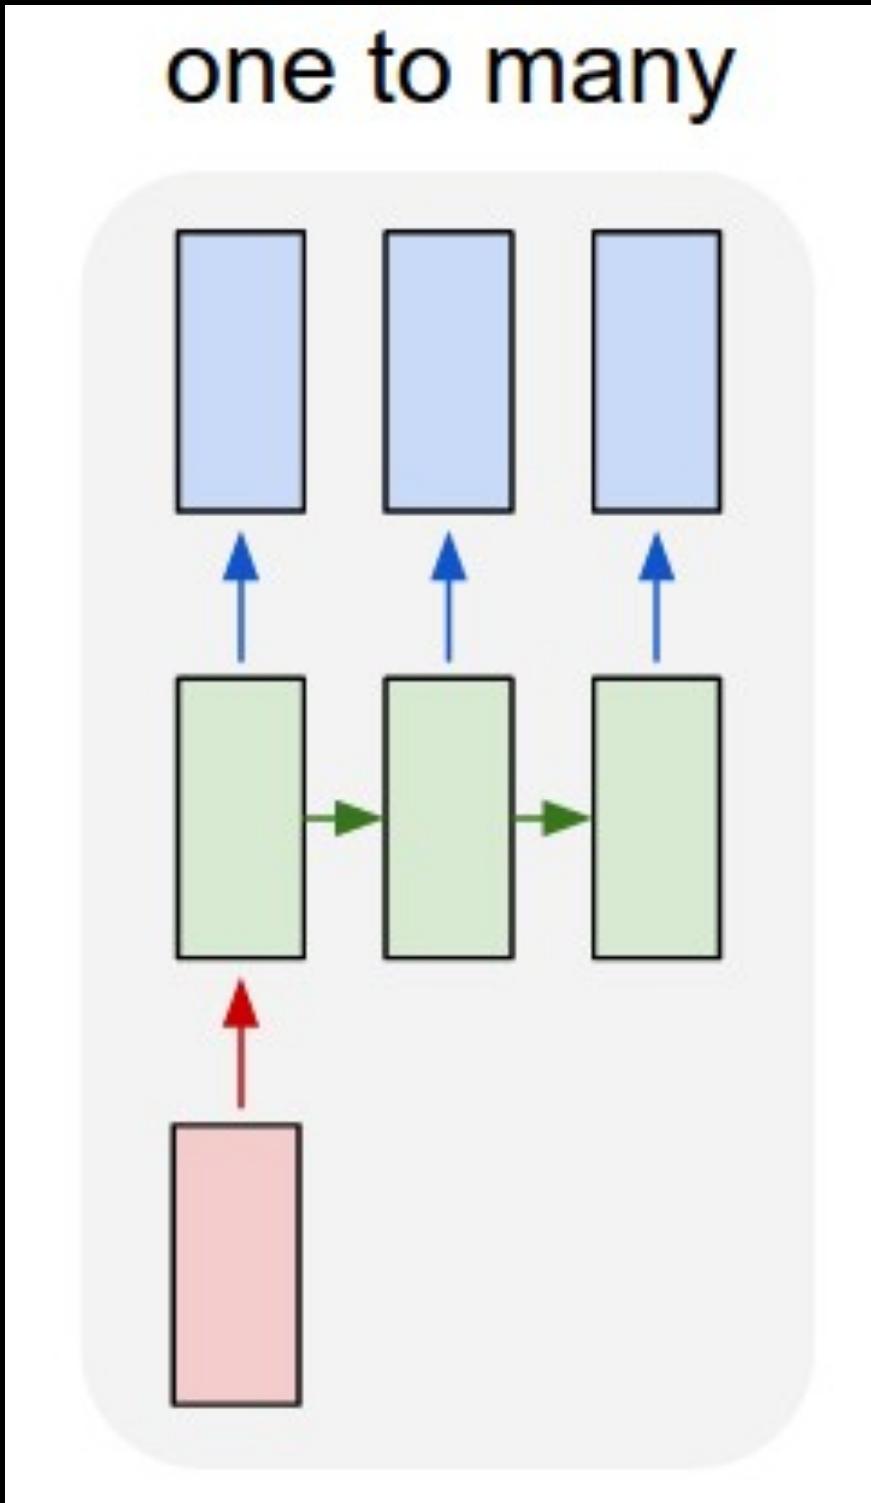
Other sequence problems



Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Other sequence problems

Image captioning



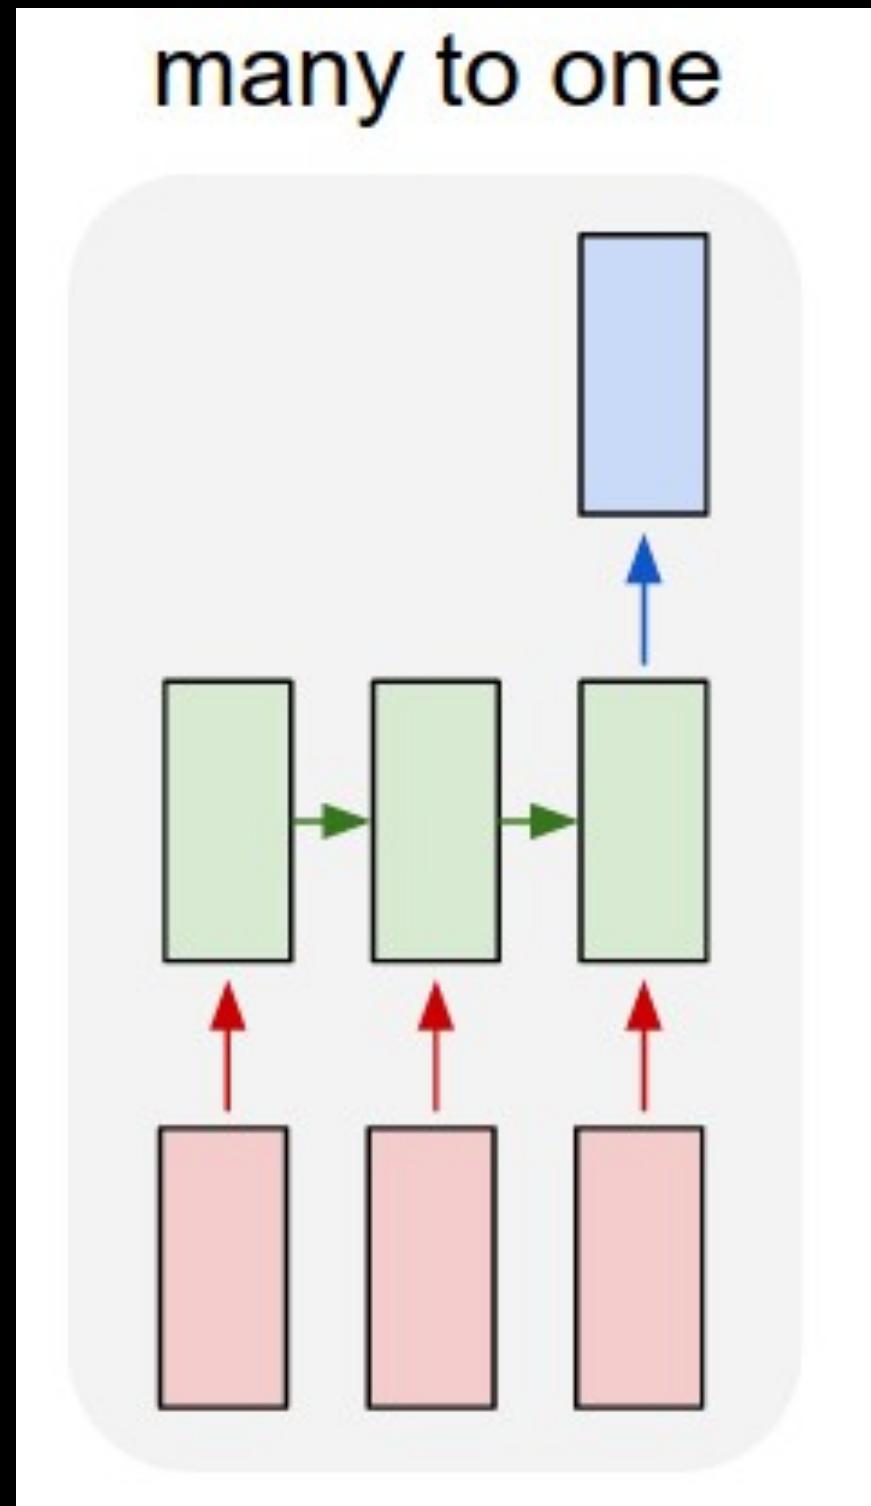
Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Input



Output A sledgehammer leaning up against a tire

Other sequence problems



Input

Output

Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Sentiment analysis

2020 Machine Learning Roadmap (still valid for 2021)

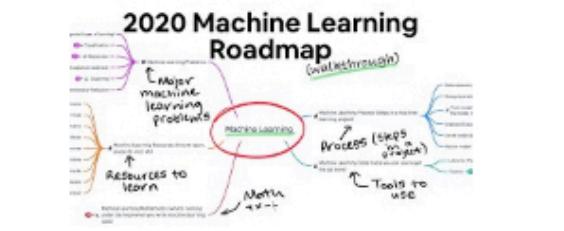
1a • 10 hours ago

content like are what makes me still believe
in humanity. thanks Daniel

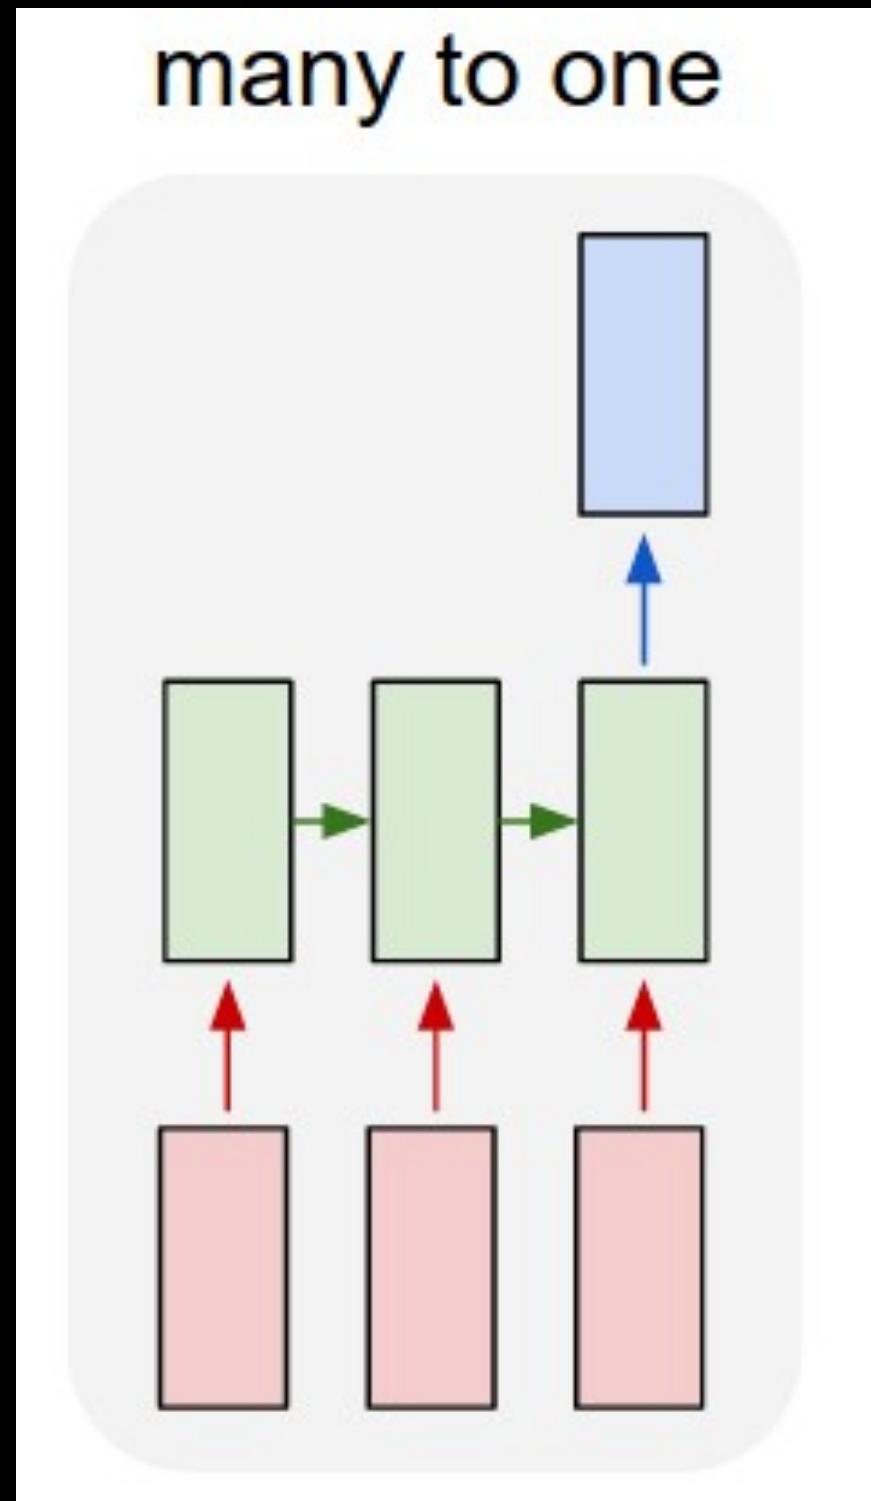


⋮

Positive



Other sequence problems



Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Input

Output

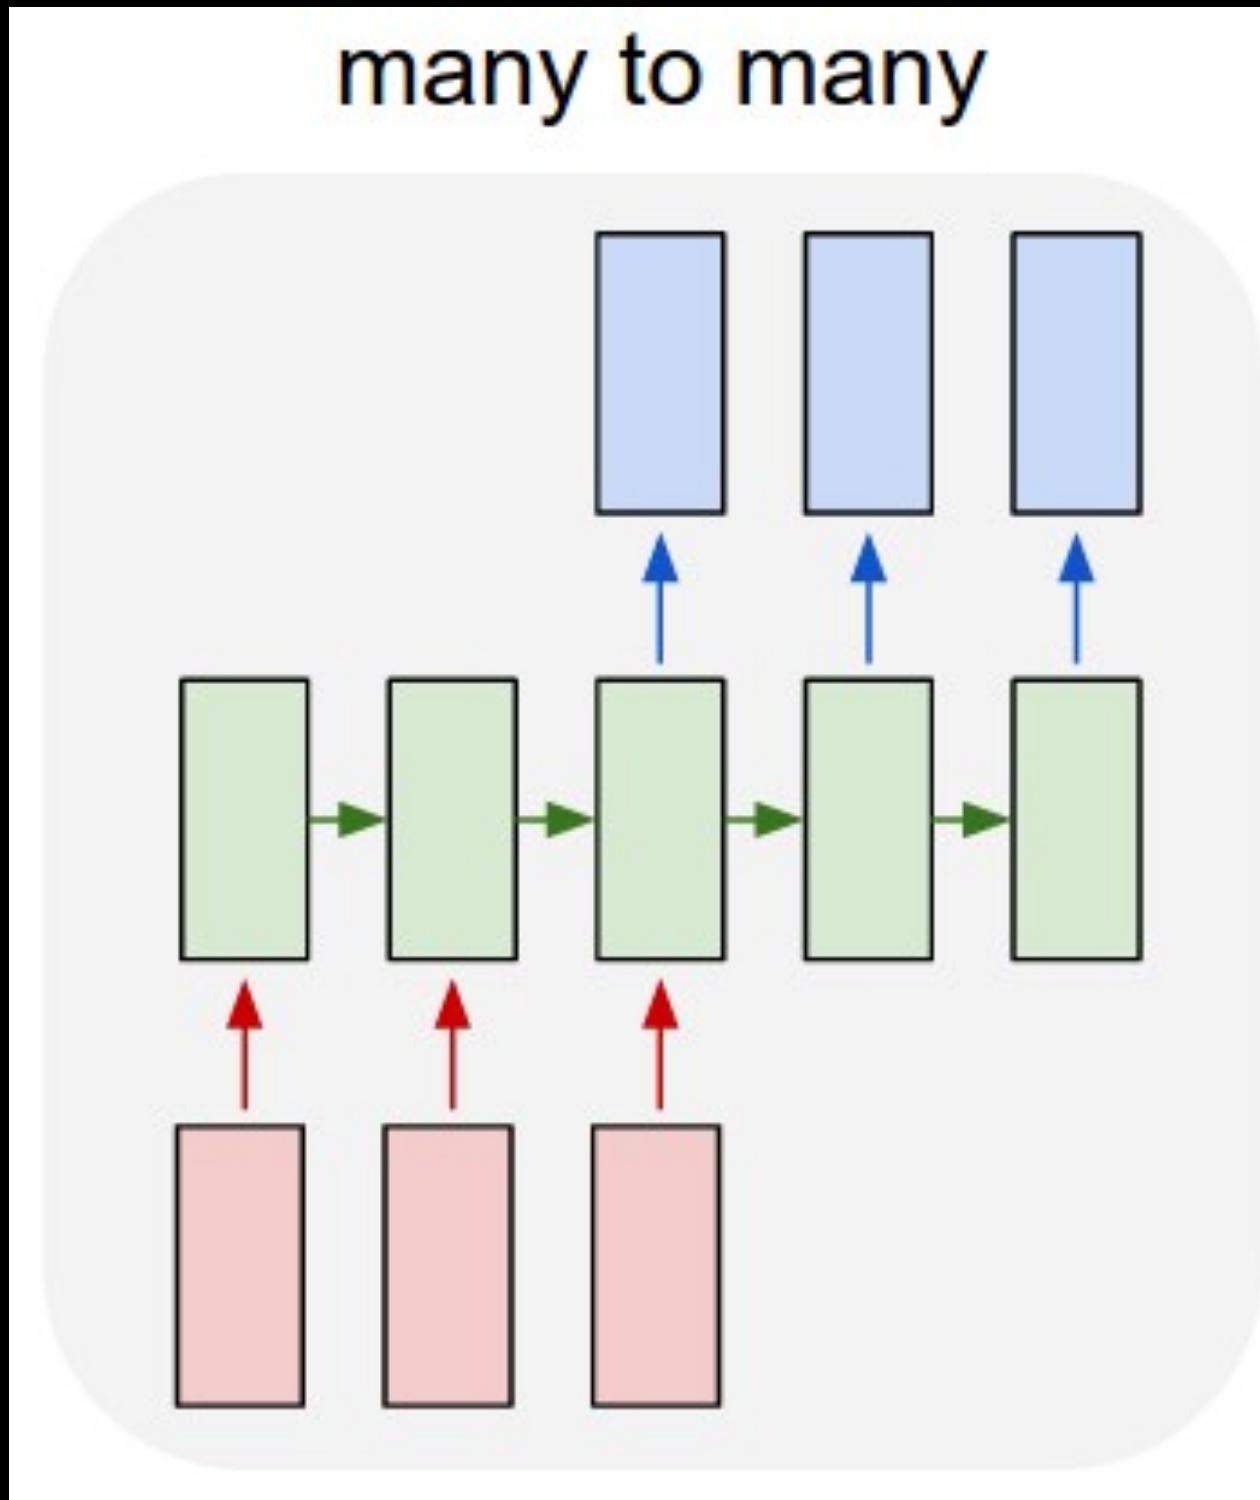


Source: <https://www.coindesk.com/price/bitcoin>

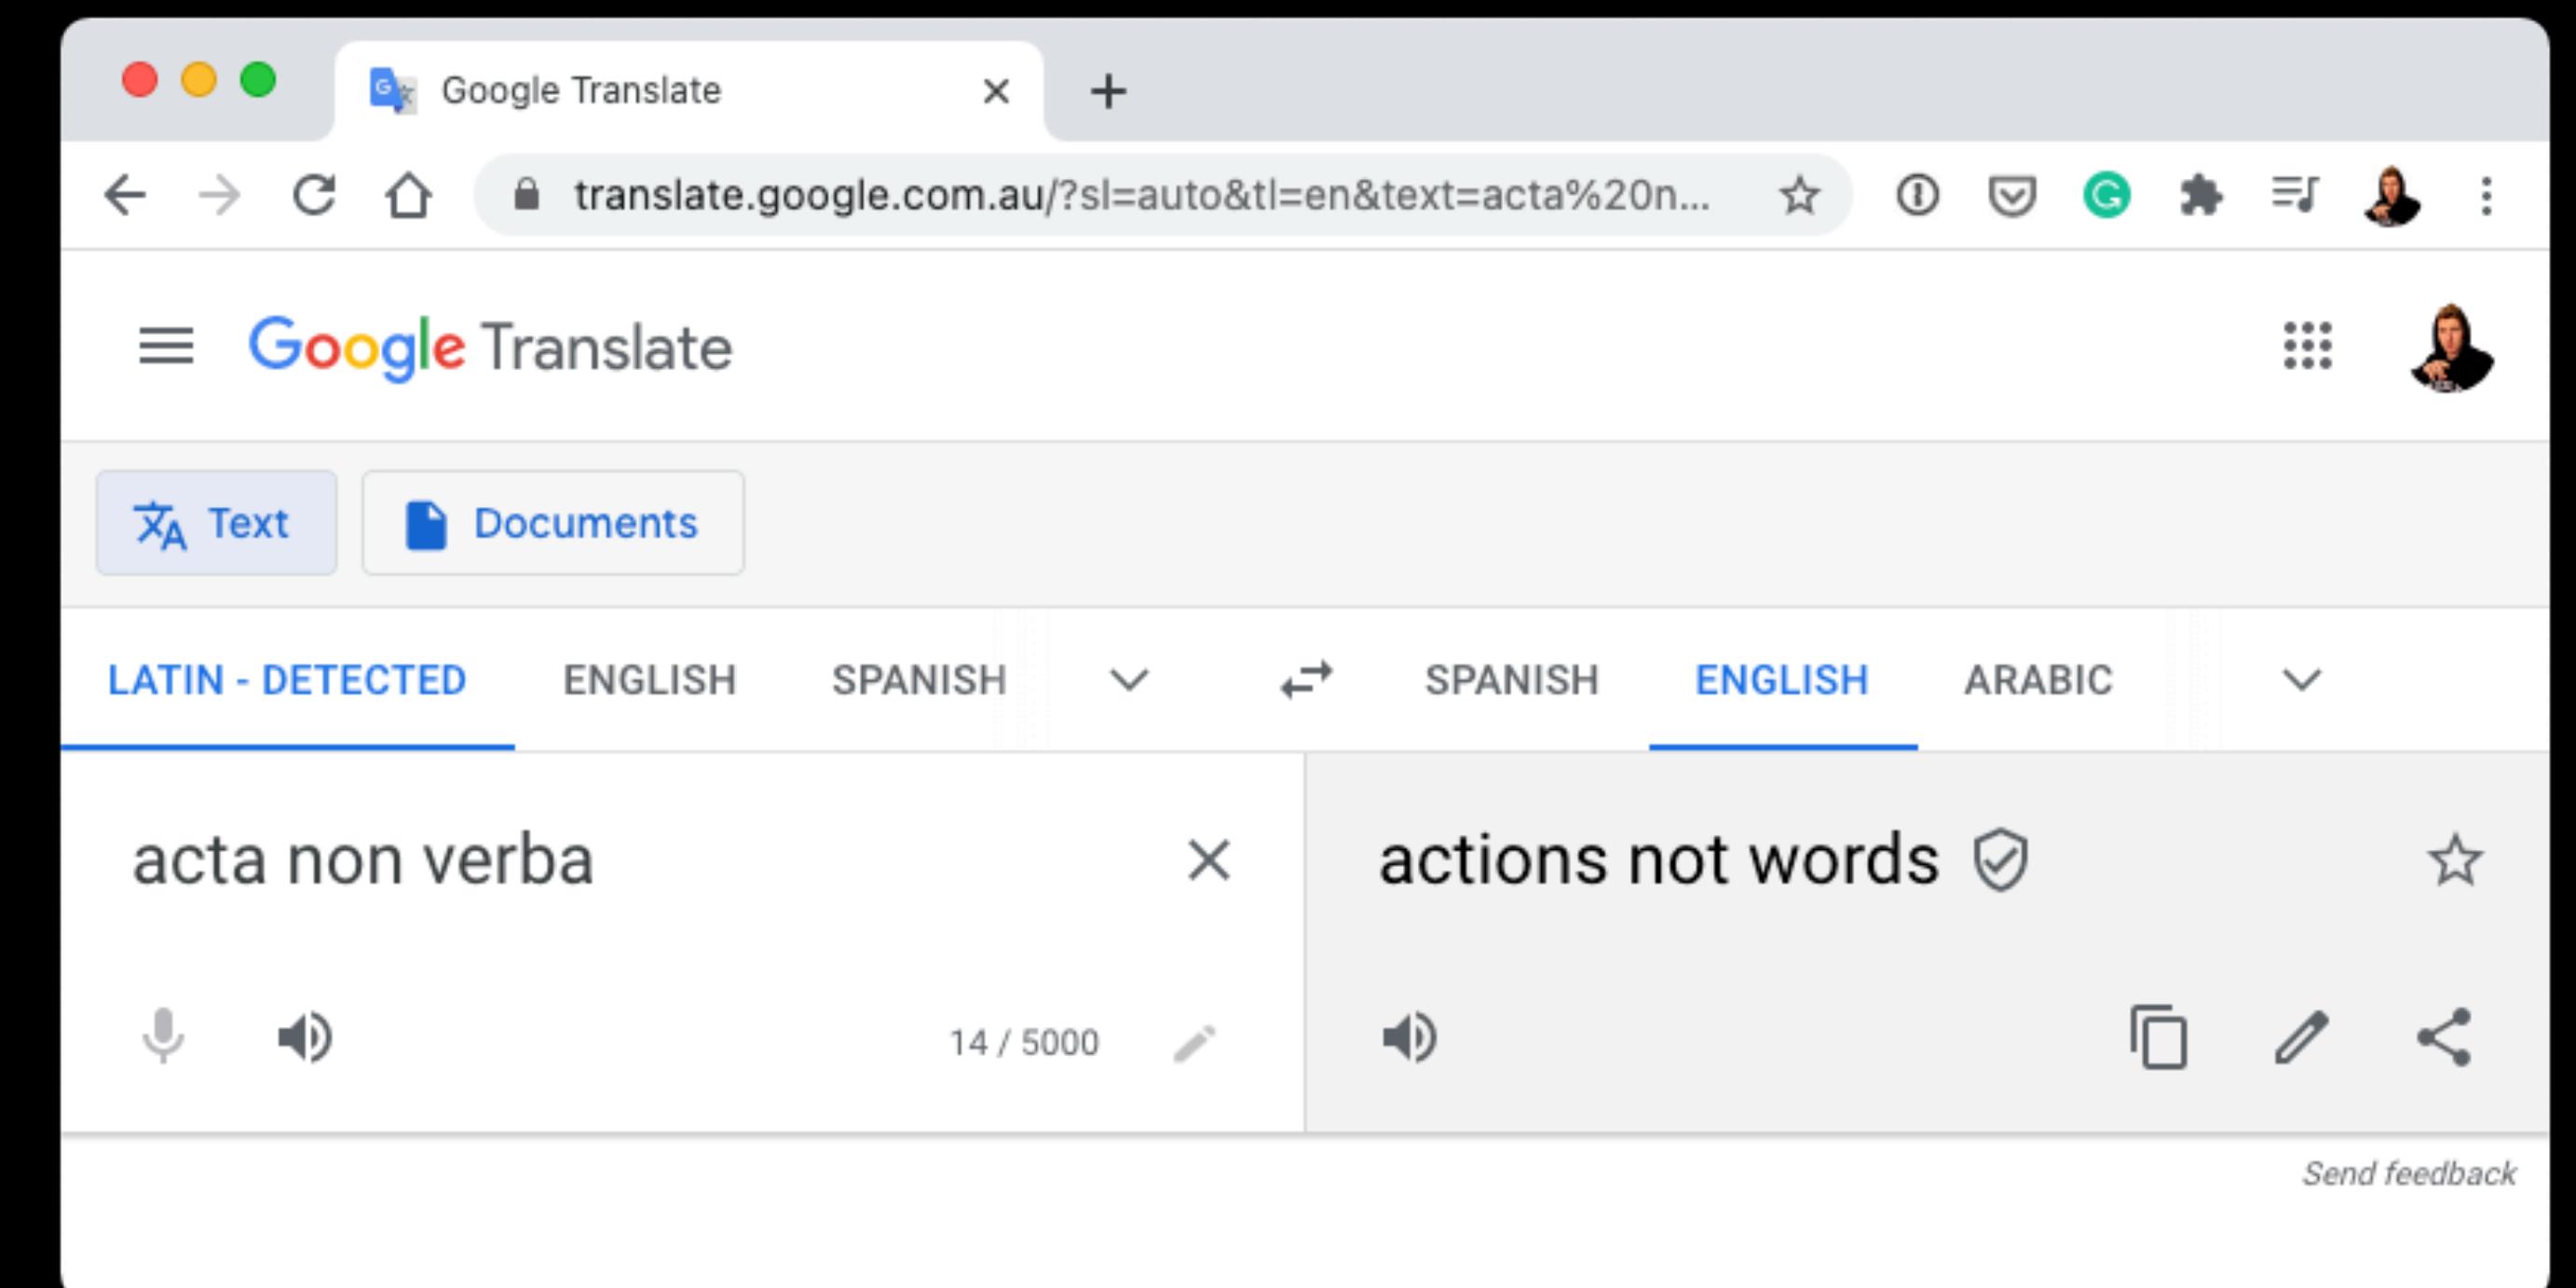
Price at next timestamp (e.g. \$59,678)

Other sequence problems

Machine Translation



Source: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>



Input

Output

What we're going to cover

(broadly)

- Downloading and preparing a text dataset
- How to prepare text data for modelling (**tokenization and embedding**)
- Setting up multiple modelling experiments with **recurrent neural networks** (RNNs)
- Building a text **feature extraction model using TensorFlow Hub**
- Finding the most wrong prediction examples
- Using a model we've built to make predictions on text from the wild

(we'll be cooking up lots of code!)

How:



NLP inputs and outputs



“Is this Tweet for a disaster or not?”



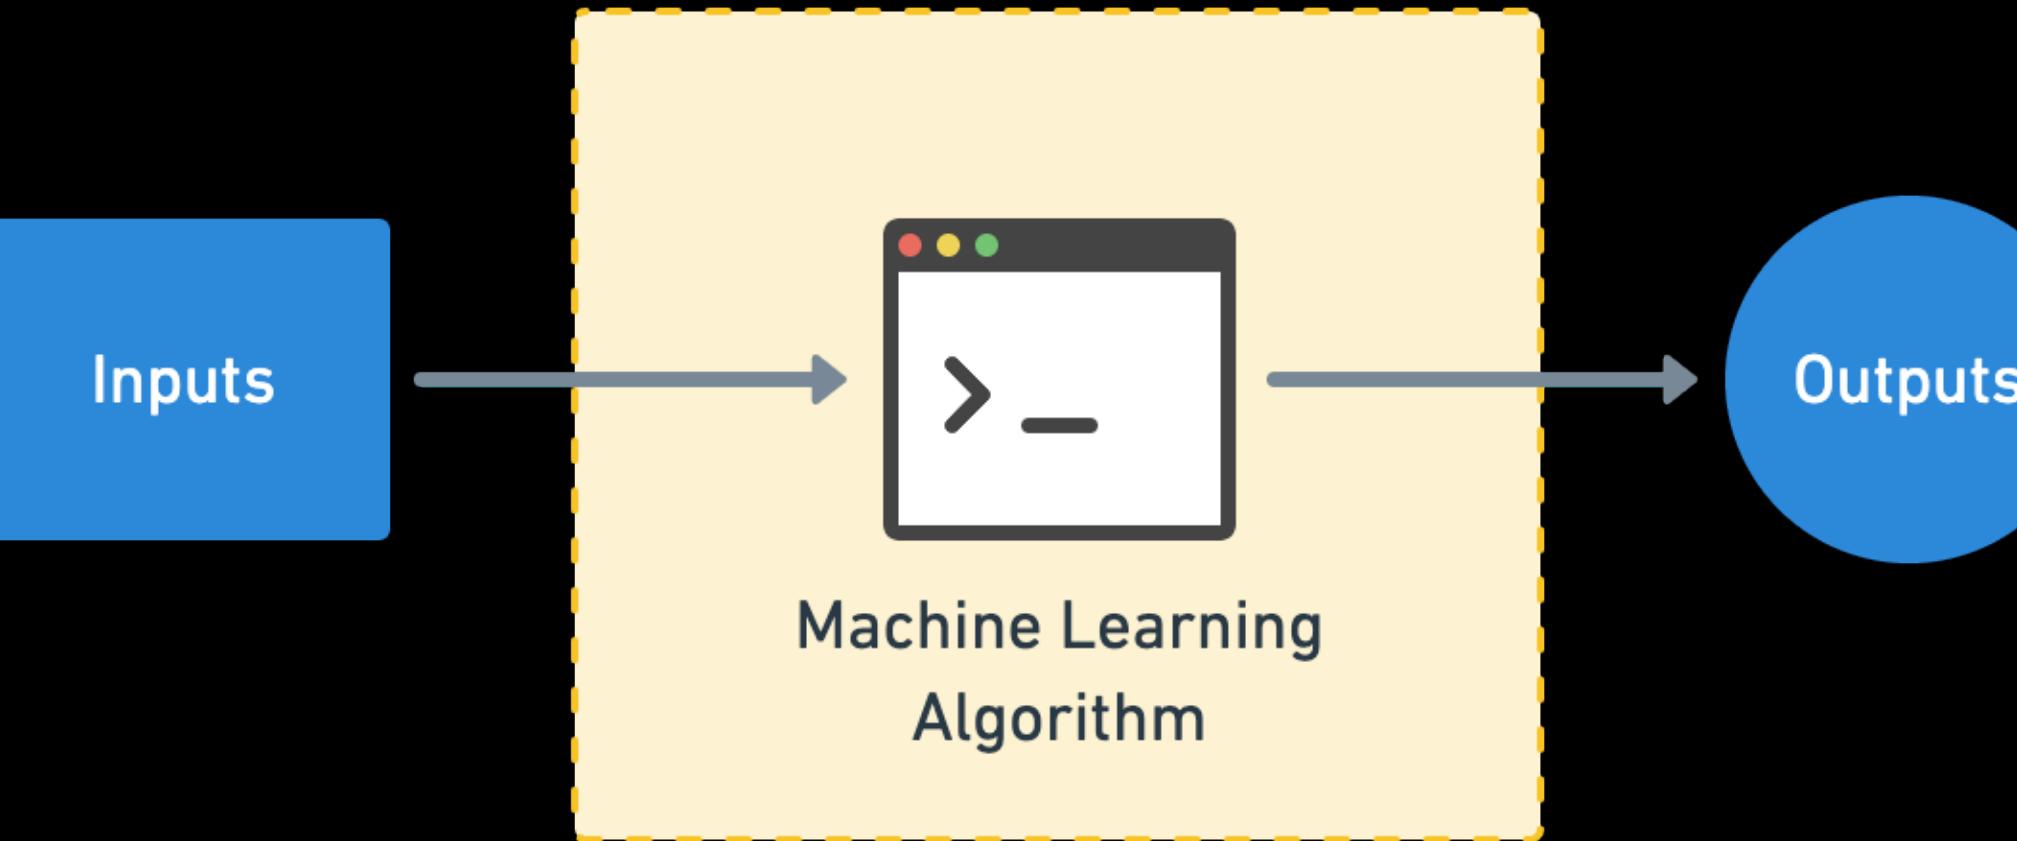
Diaster

Not Diaster

Actual output

$[[0.22, 0.98, 0.02...],$
 $[0.09, 0.55, 0.87...], \rightarrow$
 $[0.53, 0.81, 0.79...],$
 $\dots,$

Numerical encoding
(Tokenization + Embedding)



(often already exists, if not,
you can build one)

$[[0.97, 0.03],$
 $[0.81, 0.19],$
 $\dots,$

Predicted output

(comes from looking at lots
of these)

Input and output shapes

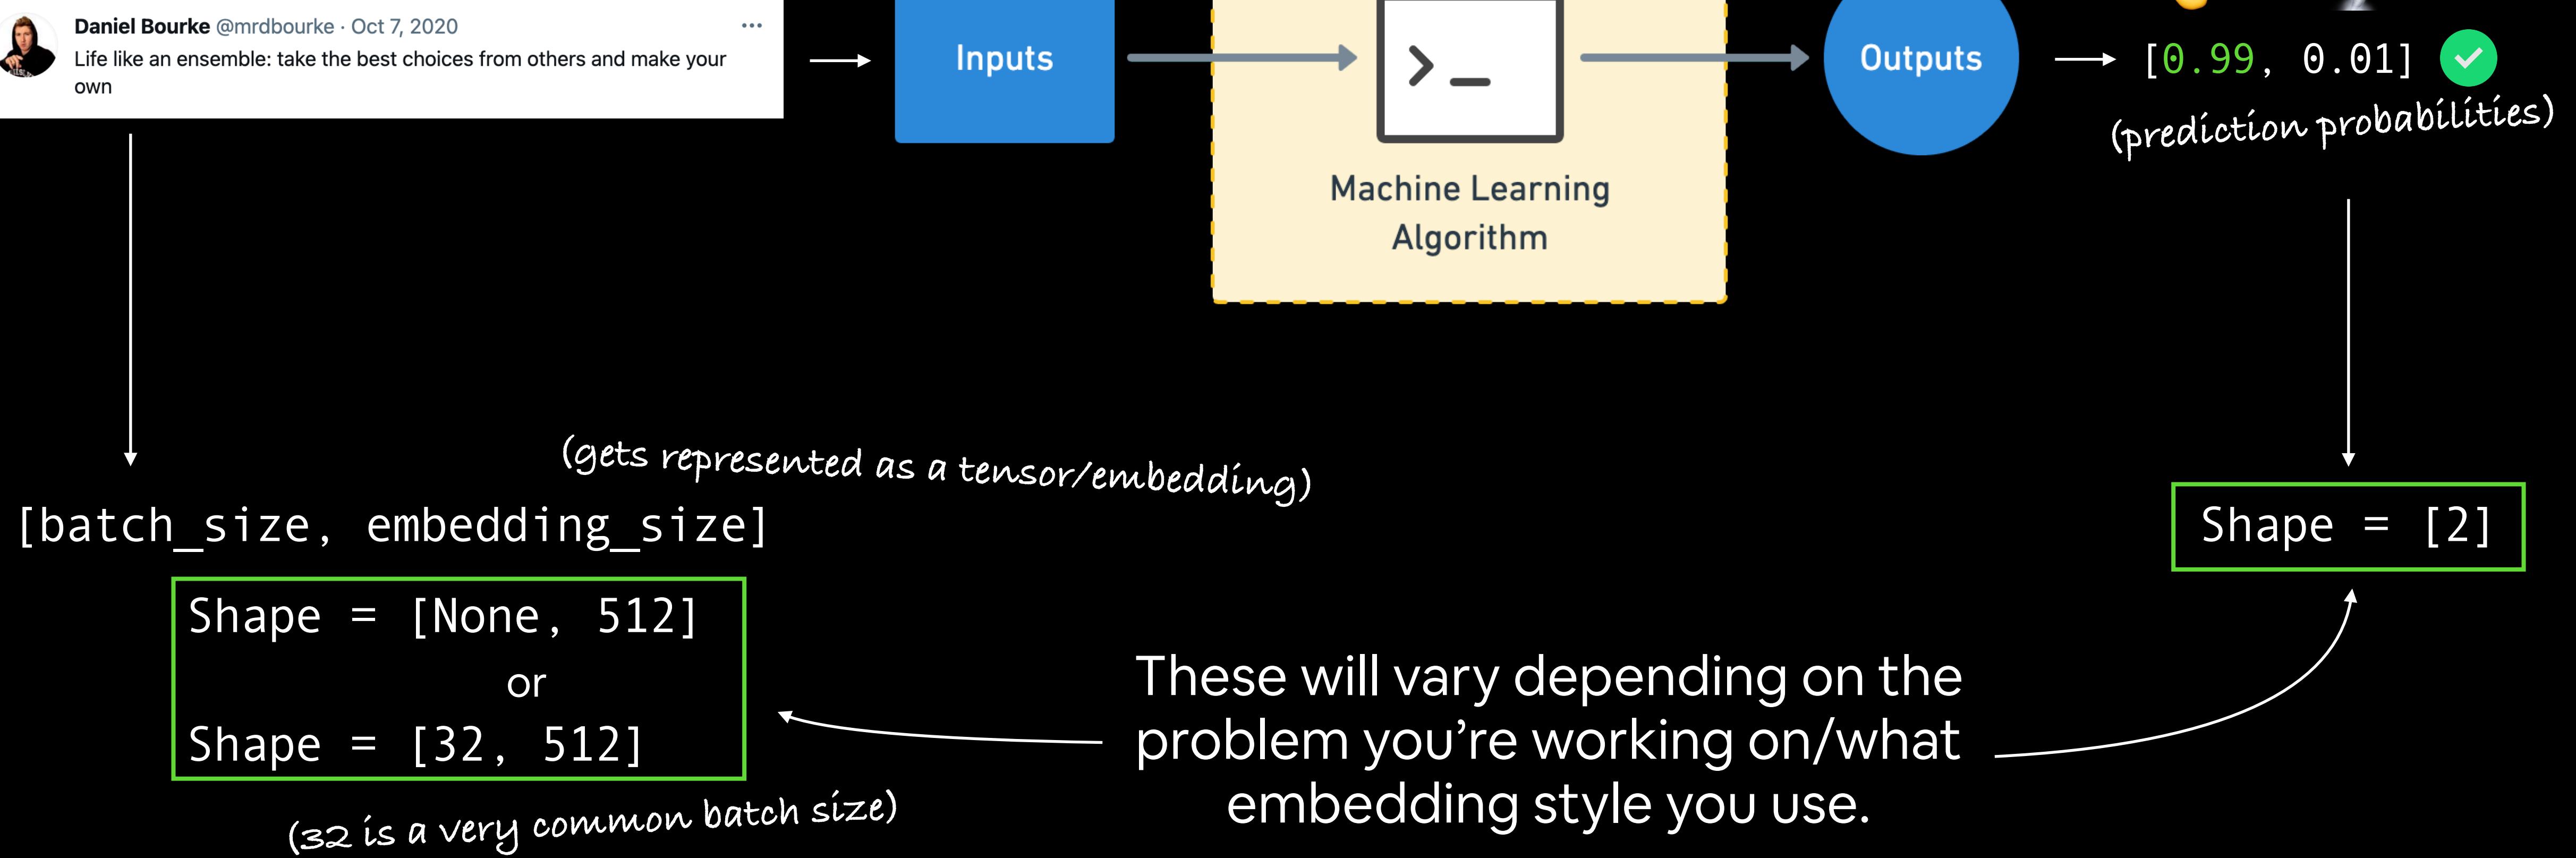
(for a text classification example)

*We're going to be building RNNs/
CNNs/Feature extractors to do this part!*

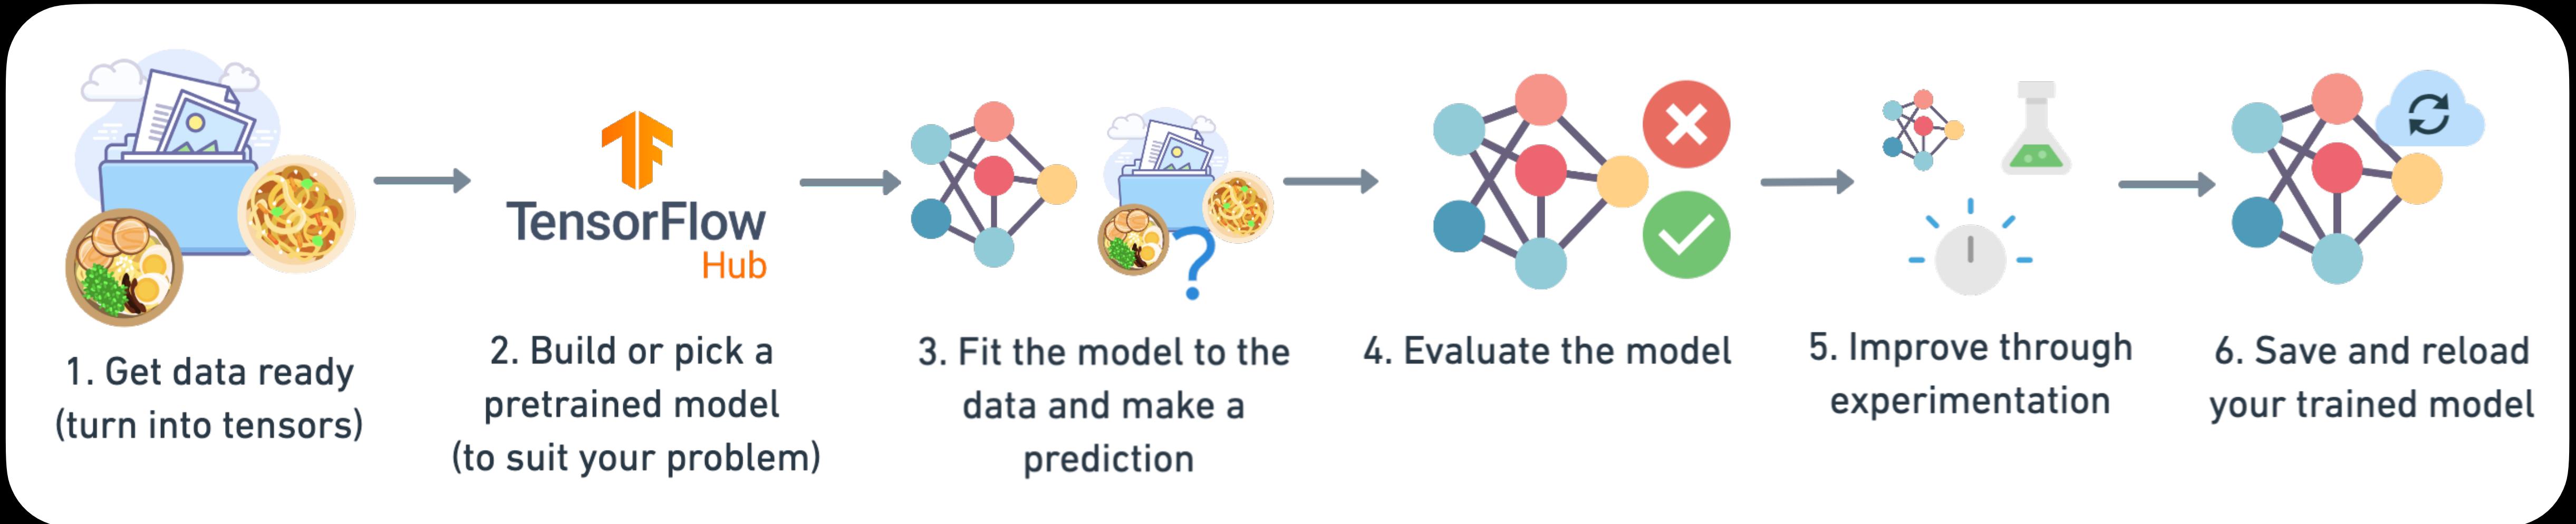


Daniel Bourke @mrdbourke · Oct 7, 2020

Life like an ensemble: take the best choices from others and make your own



Steps in modelling with TensorFlow



1. Turn all data into numbers (neural networks can't handle text/natural language)
2. Make sure all of your tensors are the right shape (pad sequences which don't fit)

“What is a recurrent neural
network (RNN)?”

(typical)*

Architecture of an RNN

Hyperparameter/Layer type	What does it do?	Typical values
Input text(s)	Target texts/sequences you'd like to discover patterns in	Whatever you can represent as text or a sequence
Input layer	Takes in target sequence	input_shape = [batch_size, embedding_size] or [batch_size, sequence_shape]
Text vectorization layer	Maps input sequences to numbers	Multiple, can create with tf.keras.layers.experimental.preprocessing.TextVectorization
Embedding	Turns mapping of text vectors to embedding matrix (representaiton of how words relate)	Multiple, can create with tf.keras.layers.Embedding
RNN cell(s)	Finds patterns in sequences	SimpleRNN , LSTM , GRU
Hidden activation	Adds non-linearity to learned features (non-straight lines)	Usually Tanh (hyperbolic tangent) (tf.keras.activations.tanh)
Pooling layer	Reduces the dimensionality of learned sequence features (usually for Conv1D models)	Average (tf.keras.layers.GlobalAveragePooling1D) or Max (tf.keras.layers.GlobalMaxPool1D)
Fully connected layer	Further refines learned features from recurrent layers	tf.keras.layers.Dense
Output layer	Takes learned features and outputs them in shape of target labels	output_shape = [number_of_classes] (e.g. 2 for Diaster, Not Diaster)
Output activation	Adds non-linearities to output layer	tf.keras.activations.sigmoid (binary classification) or tf.keras.activations.softmax

```
# 1. Create LSTM model
from tensorflow.keras import layers
inputs = layers.Input(shape=(1,), dtype="string")
x = text_vectorizer(inputs) # turn input sequence to numbers
x = embedding(x) # create embedding matrix
x = layers.LSTM(64, activation="tanh")(x) # return vector for whole sequence
outputs = layers.Dense(1, activation="sigmoid")(x) # create output layer
model = tf.keras.Model(inputs, outputs, name="LSTM_model")

# 2. Compile model
model.compile(loss=tf.keras.losses.BinaryCrossentropy(),
              optimizer=tf.keras.optimizers.Adam(),
              metrics=["accuracy"])

# 3. Fit the model
history = model.fit(train_sentences, train_labels, epochs=5)
```

(what we're working towards building)

→ Not Diaster ✌

Daniel Bourke @mrdbourke · Oct 7, 2020
Life like an ensemble: take the best choices from others and make your own

...



*Note: there are almost an unlimited amount of ways you could stack together a recurrent neural network, this slide demonstrates only one.

Let's code!

Tokenization vs Embedding

I love TensorFlow

Tokenization — straight mapping from token to number (can be modelled but quickly gets too big)

Embedding — richer representation of relationships between tokens (can limit size + can be learned)



$l = 0$
love = 1
TensorFlow = 2

[[1, 0, 0],
[0, 1, 0], One-hot
[0, 0, 1], Encoding

...,
[[0.492, 0.005, 0.019],
[0.060, 0.233, 0.899], Embedding
[0.741, 0.983, 0.567],
...,

Experiments we're running

Experiment Number	Model
0	Naive Bayes with TF-IDF encoder (baseline)
1	Feed-forward neural network (dense model)
2	LSTM (RNN)
3	GRU (RNN)
4	Bidirectional-LSTM (RNN)
5	1D Convolutional Neural Network
6	TensorFlow Hub Pretrained Feature Extractor
7	TensorFlow Hub Pretrained Feature Extractor (10% of data)

(some common)

Classification evaluation methods

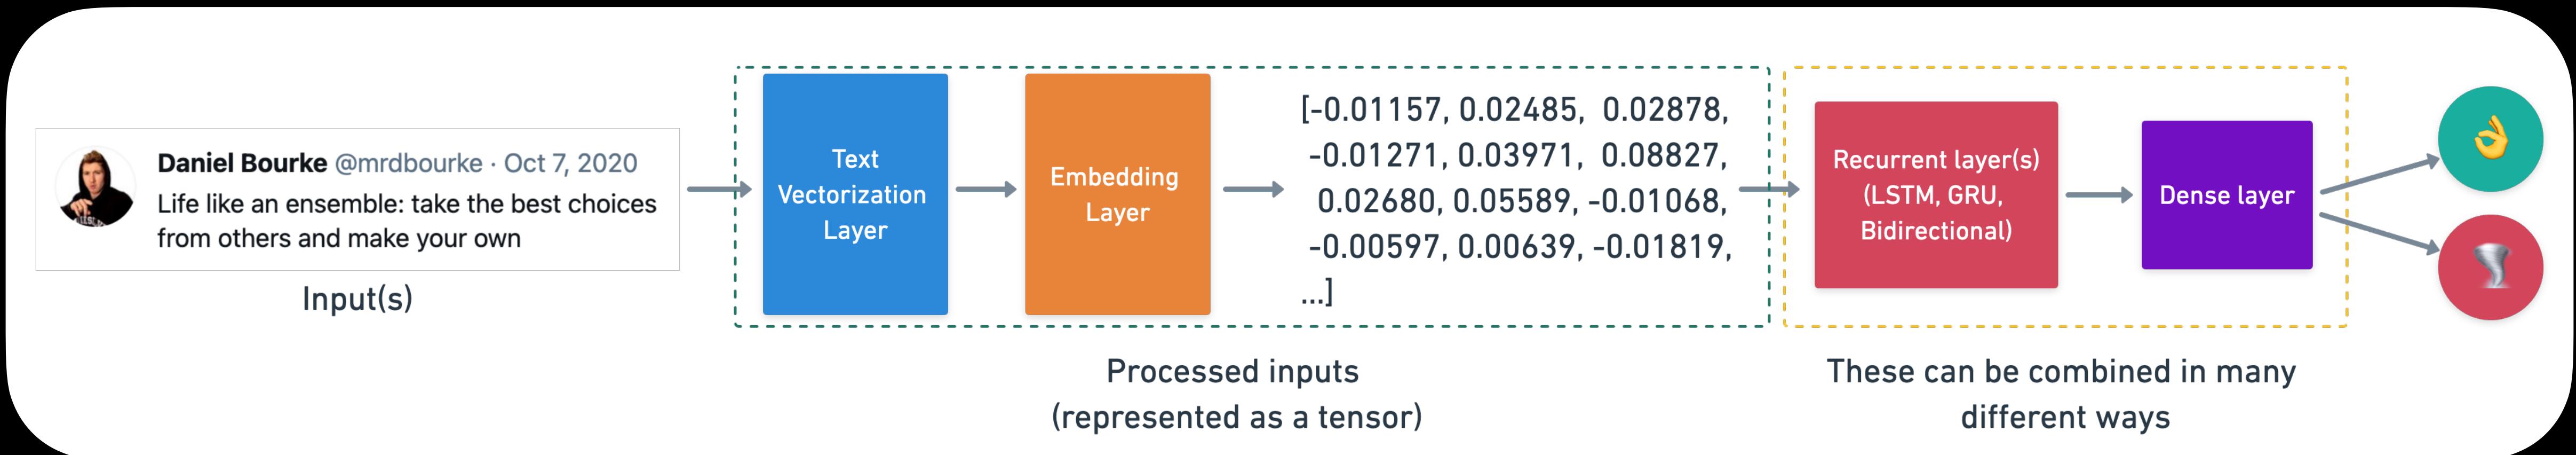
Key: **tp** = True Positive, **tn** = True Negative, **fp** = False Positive, **fn** = False Negative

Metric Name	Metric Forumla	Code	When to use
Accuracy	Accuracy = $\frac{tp + tn}{tp + tn + fp + fn}$	<code>tf.keras.metrics.Accuracy()</code> or <code>sklearn.metrics.accuracy_score()</code>	Default metric for classification problems. Not the best for imbalanced classes.
Precision	Precision = $\frac{tp}{tp + fp}$	<code>tf.keras.metrics.Precision()</code> or <code>sklearn.metrics.precision_score()</code>	Higher precision leads to less false positives.
Recall	Recall = $\frac{tp}{tp + fn}$	<code>tf.keras.metrics.Recall()</code> or <code>sklearn.metrics.recall_score()</code>	Higher recall leads to less false negatives.
F1-score	F1-score = $2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$	<code>sklearn.metrics.f1_score()</code>	Combination of precision and recall, usually a good overall metric for a classification model.
Confusion matrix	NA	Custom function or <code>sklearn.metrics.confusion_matrix()</code>	When comparing predictions to truth labels to see where model gets confused. Can be hard to use with large numbers of classes.

Architecture of a RNN

(coloured block edition)

Standard RNN



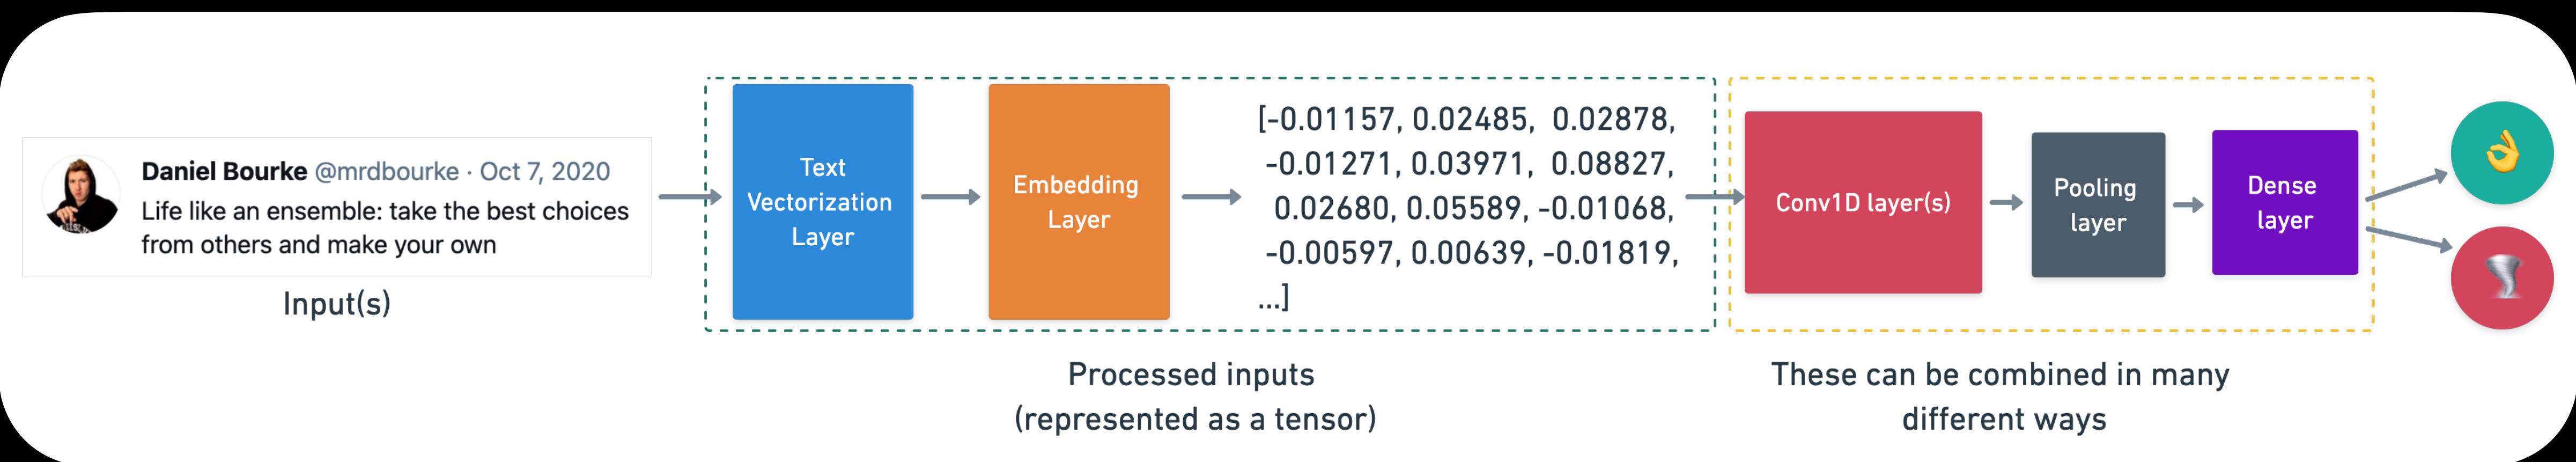
Types of RNN cells

Name	When to use	Learn more	Code
LSTM (long short-term memory)	Default RNN layer for sequence problems.	<u>Understanding LSTM Networks by Chris Olah</u>	<code>tf.keras.layers.LSTM</code>
GRU (gated recurrent unit)	Performs very similar to LSTM (could be used as a default).	<u>Illustrated Guide to LSTM's and GRU's by Michael Phi</u>	<code>tf.keras.layers.GRU</code>
Bidirectional LSTM (goes forward and backwards on sequence)	Good for sequences which may benefit from passing forward and backwards (e.g. translation or longer passages of text).	Same as above	<code>tf.keras.layers.Bidirectional</code>

Architecture of a Sequence Conv1D Model

(coloured block edition)

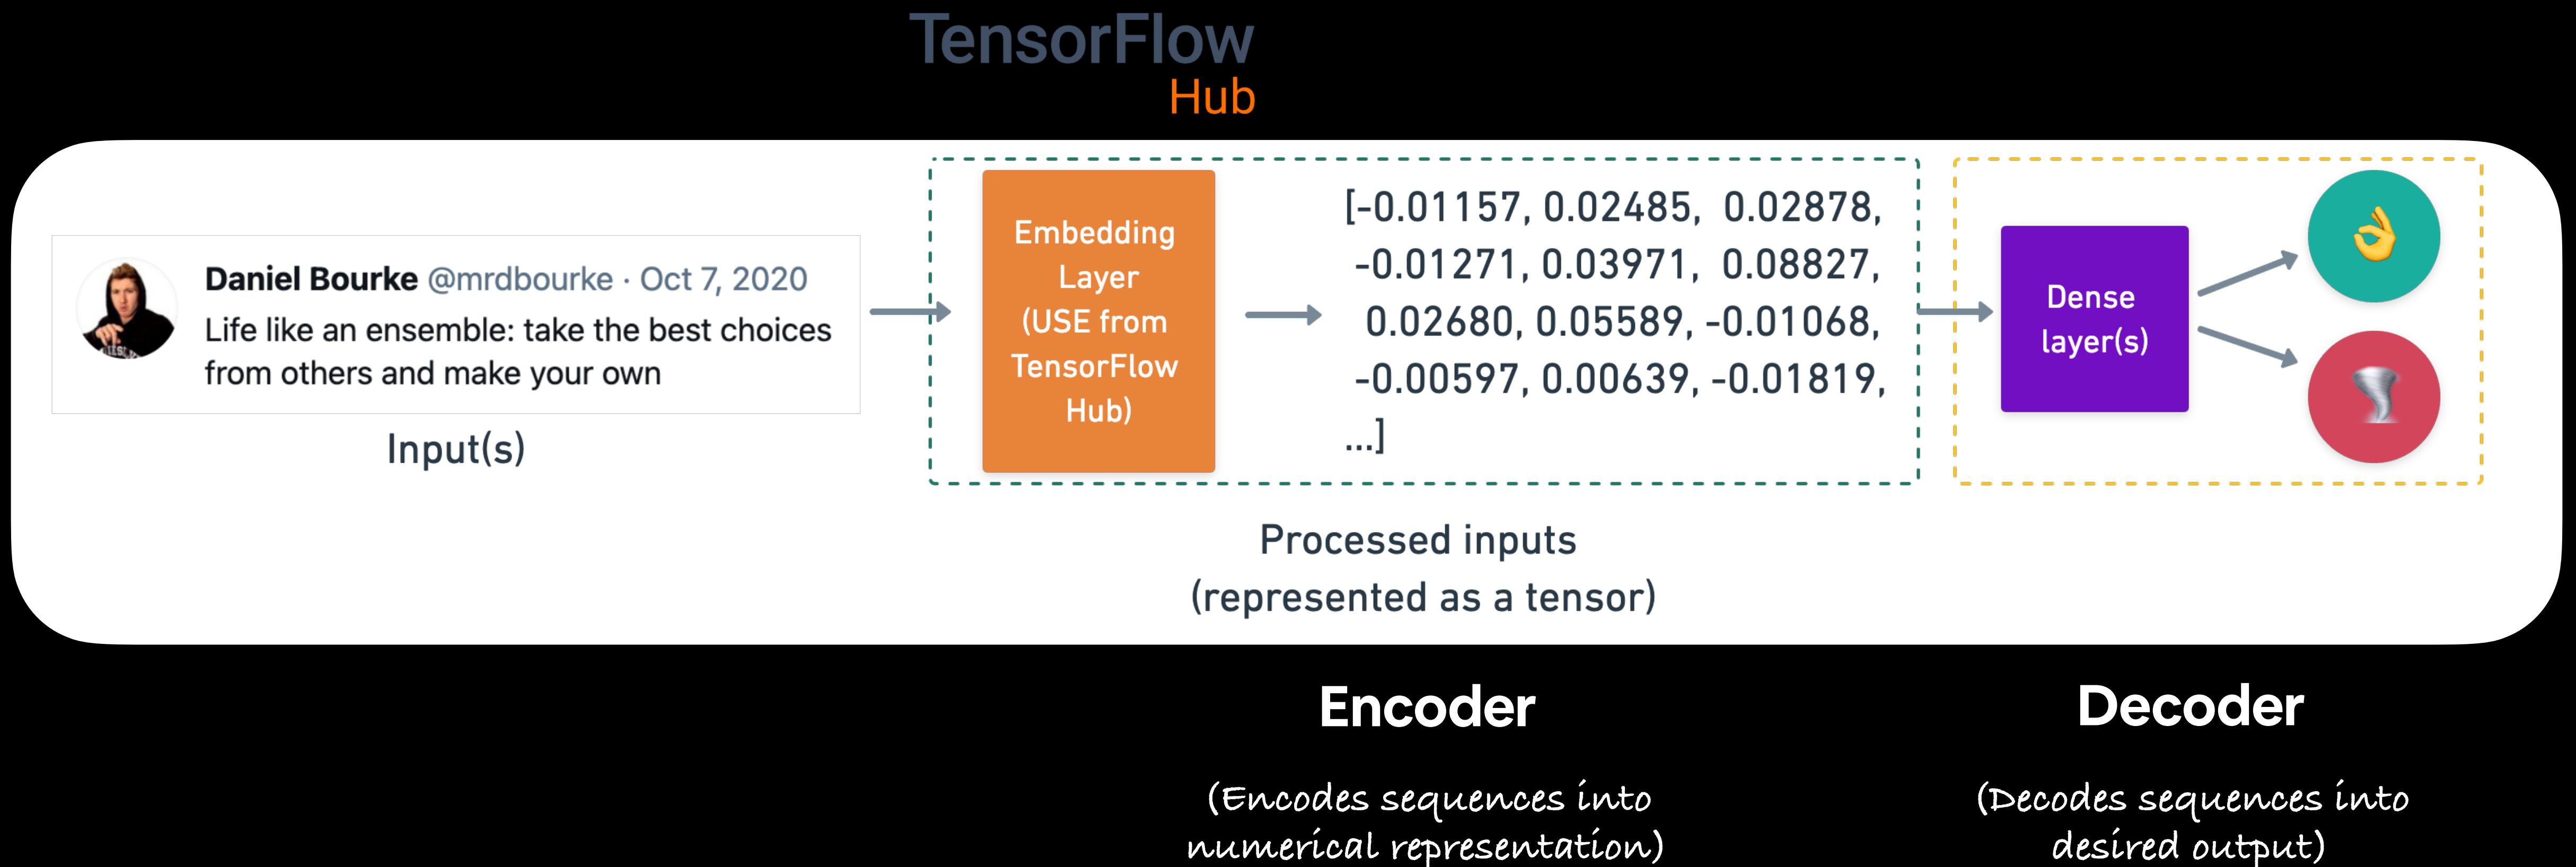
Conv1D Sequence model



Model we're building (USE* feature extractor)

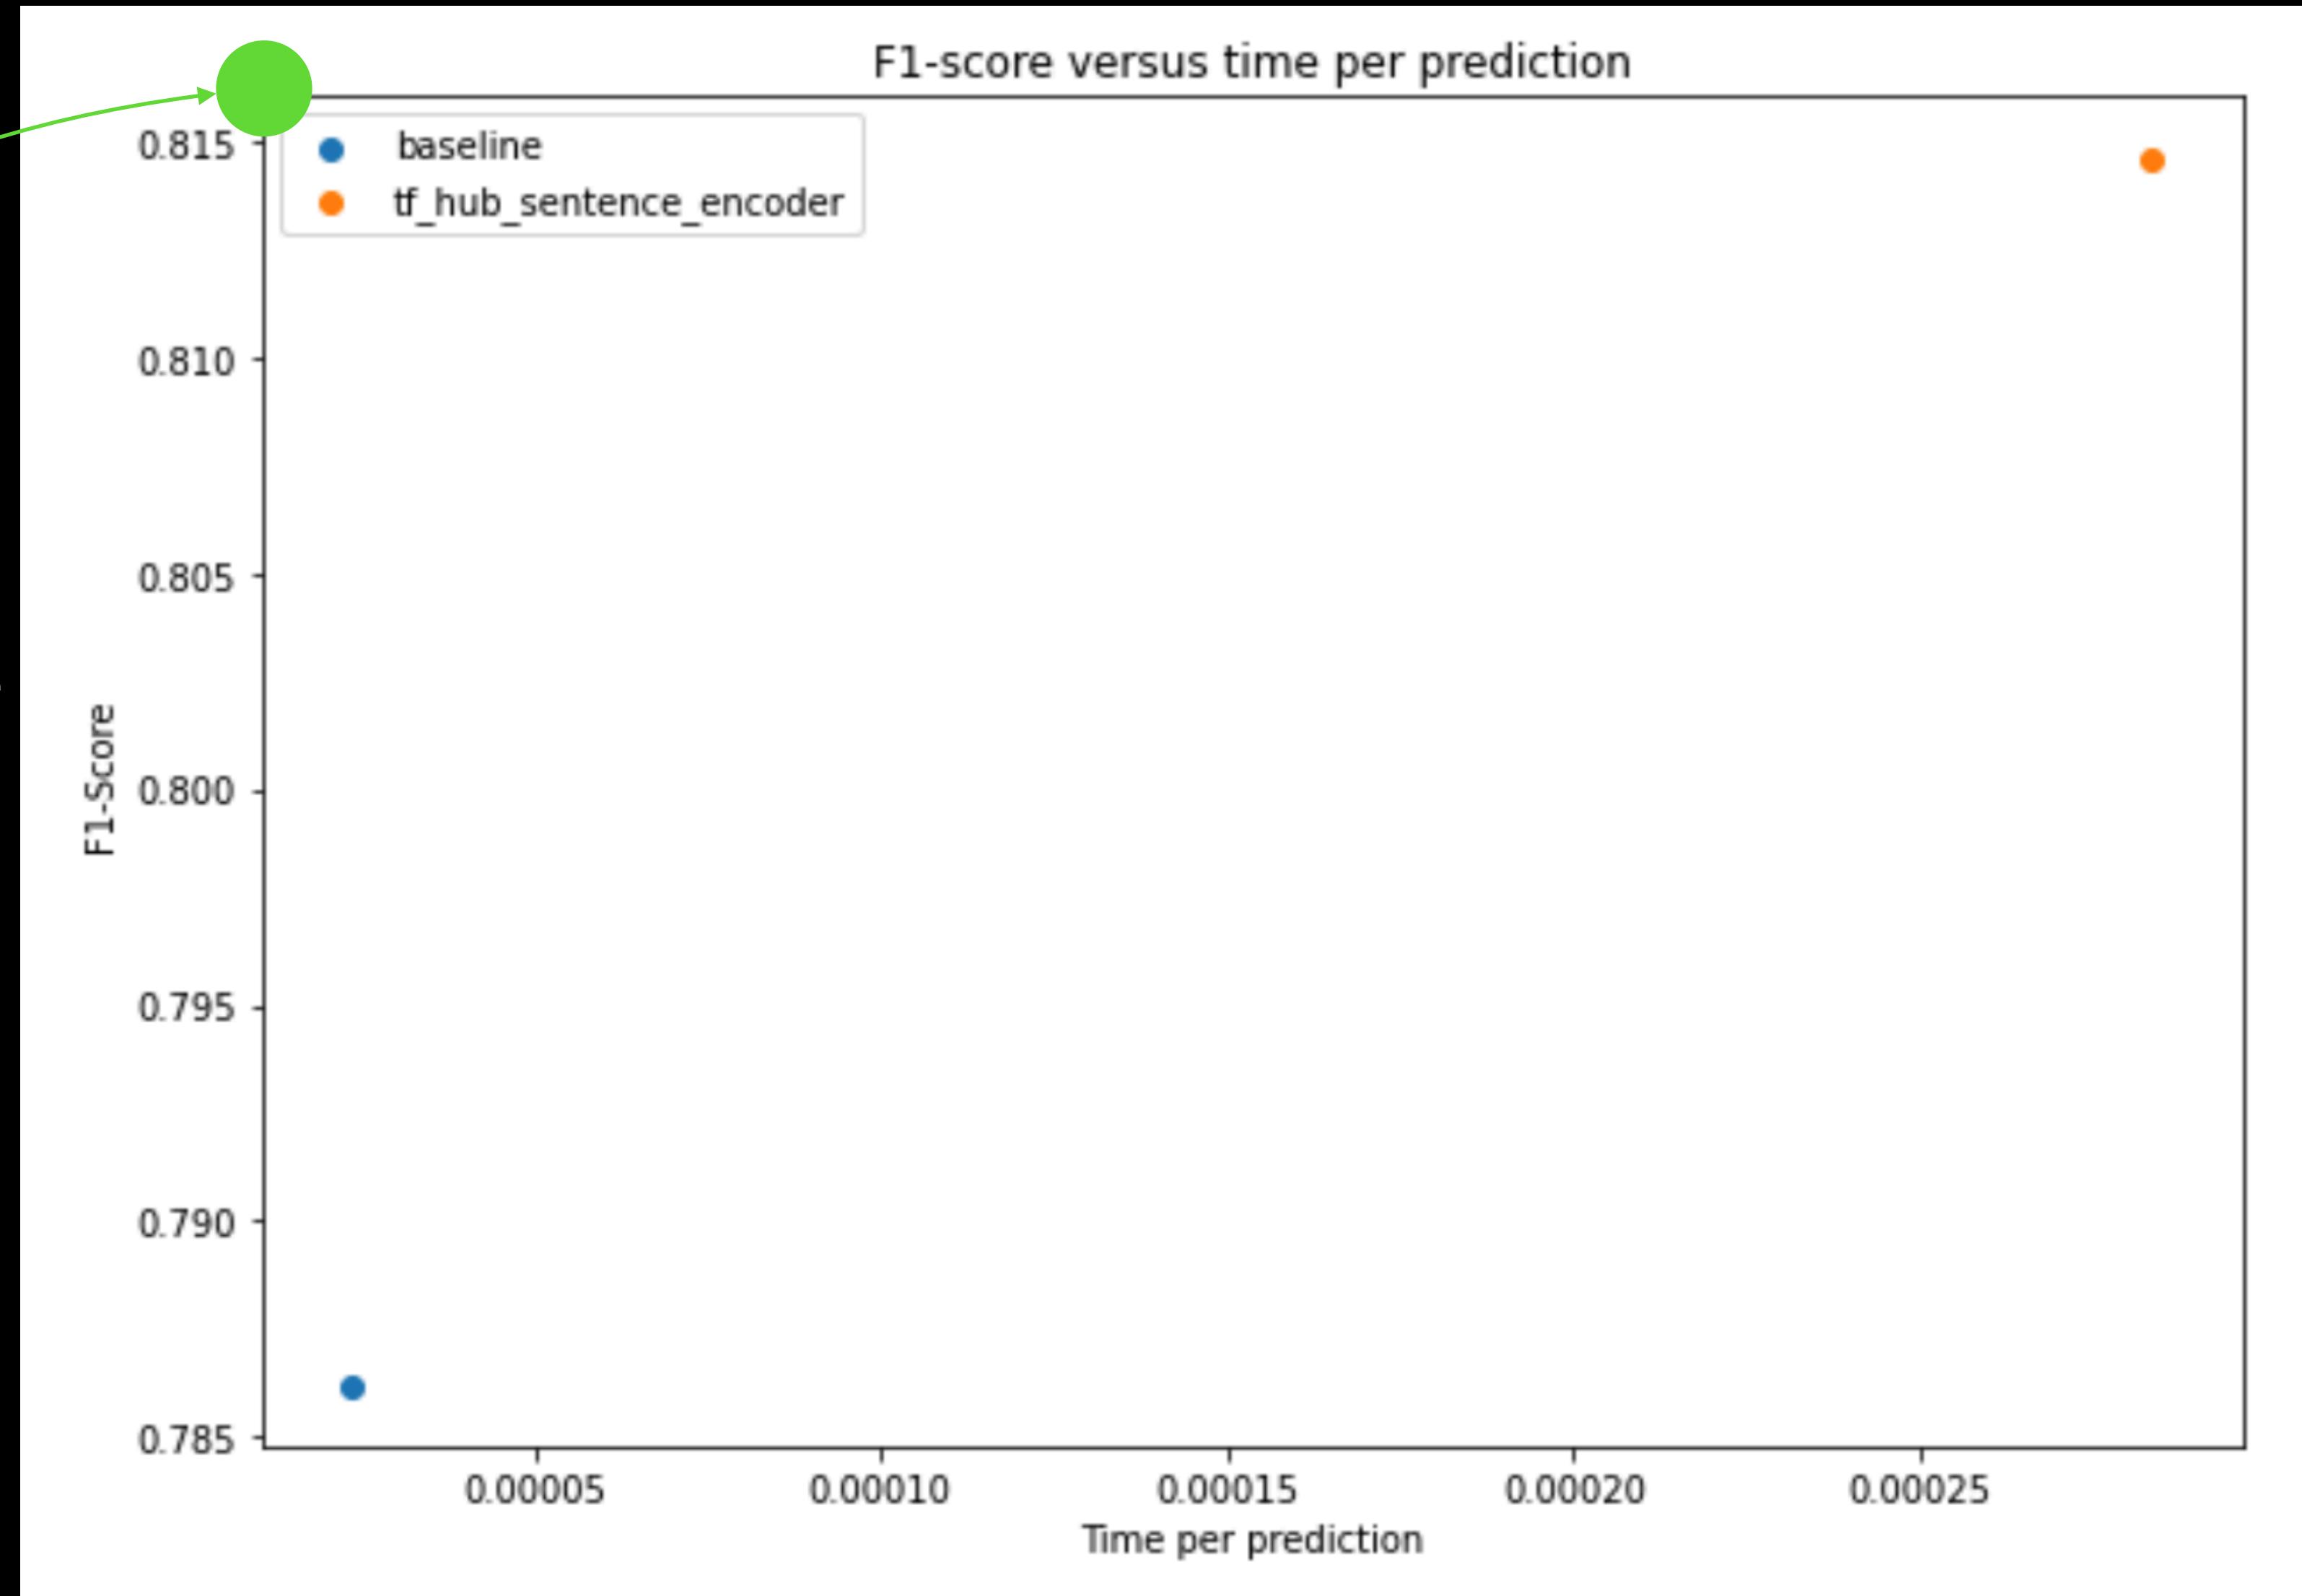
*USE = *universal Sentence Encoder*

Source: <https://tfhub.dev/google/universal-sentence-encoder/4>



Ideal speed/performance trade off

Ideal position for
speed/performance
(high performance +
high speed)



Improving a model

(from a model's perspective)

```
# 1. Create the model (specified to your problem)
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(4, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])

# 2. Compile the model
model.compile(loss=tf.keras.losses.BinaryCrossentropy(),
              optimizer=tf.keras.optimizers.Adam(lr=0.001),
              metrics=[ "accuracy" ])

# 3. Fit the model
model.fit(x_train_subset, y_train_subset, epochs=5)
```

Smaller model

```
# 1. Create the model (specified to your problem)
model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(10, activation="softmax")
])

# 2. Compile the model
model.compile(loss=tf.keras.losses.BinaryCrossentropy(),
              optimizer=tf.keras.optimizers.Adam(lr=0.0001),
              metrics=[ "accuracy" ])

# 3. Fit the model
model.fit(x_train_full, y_train_full, epochs=100)
```

Larger model

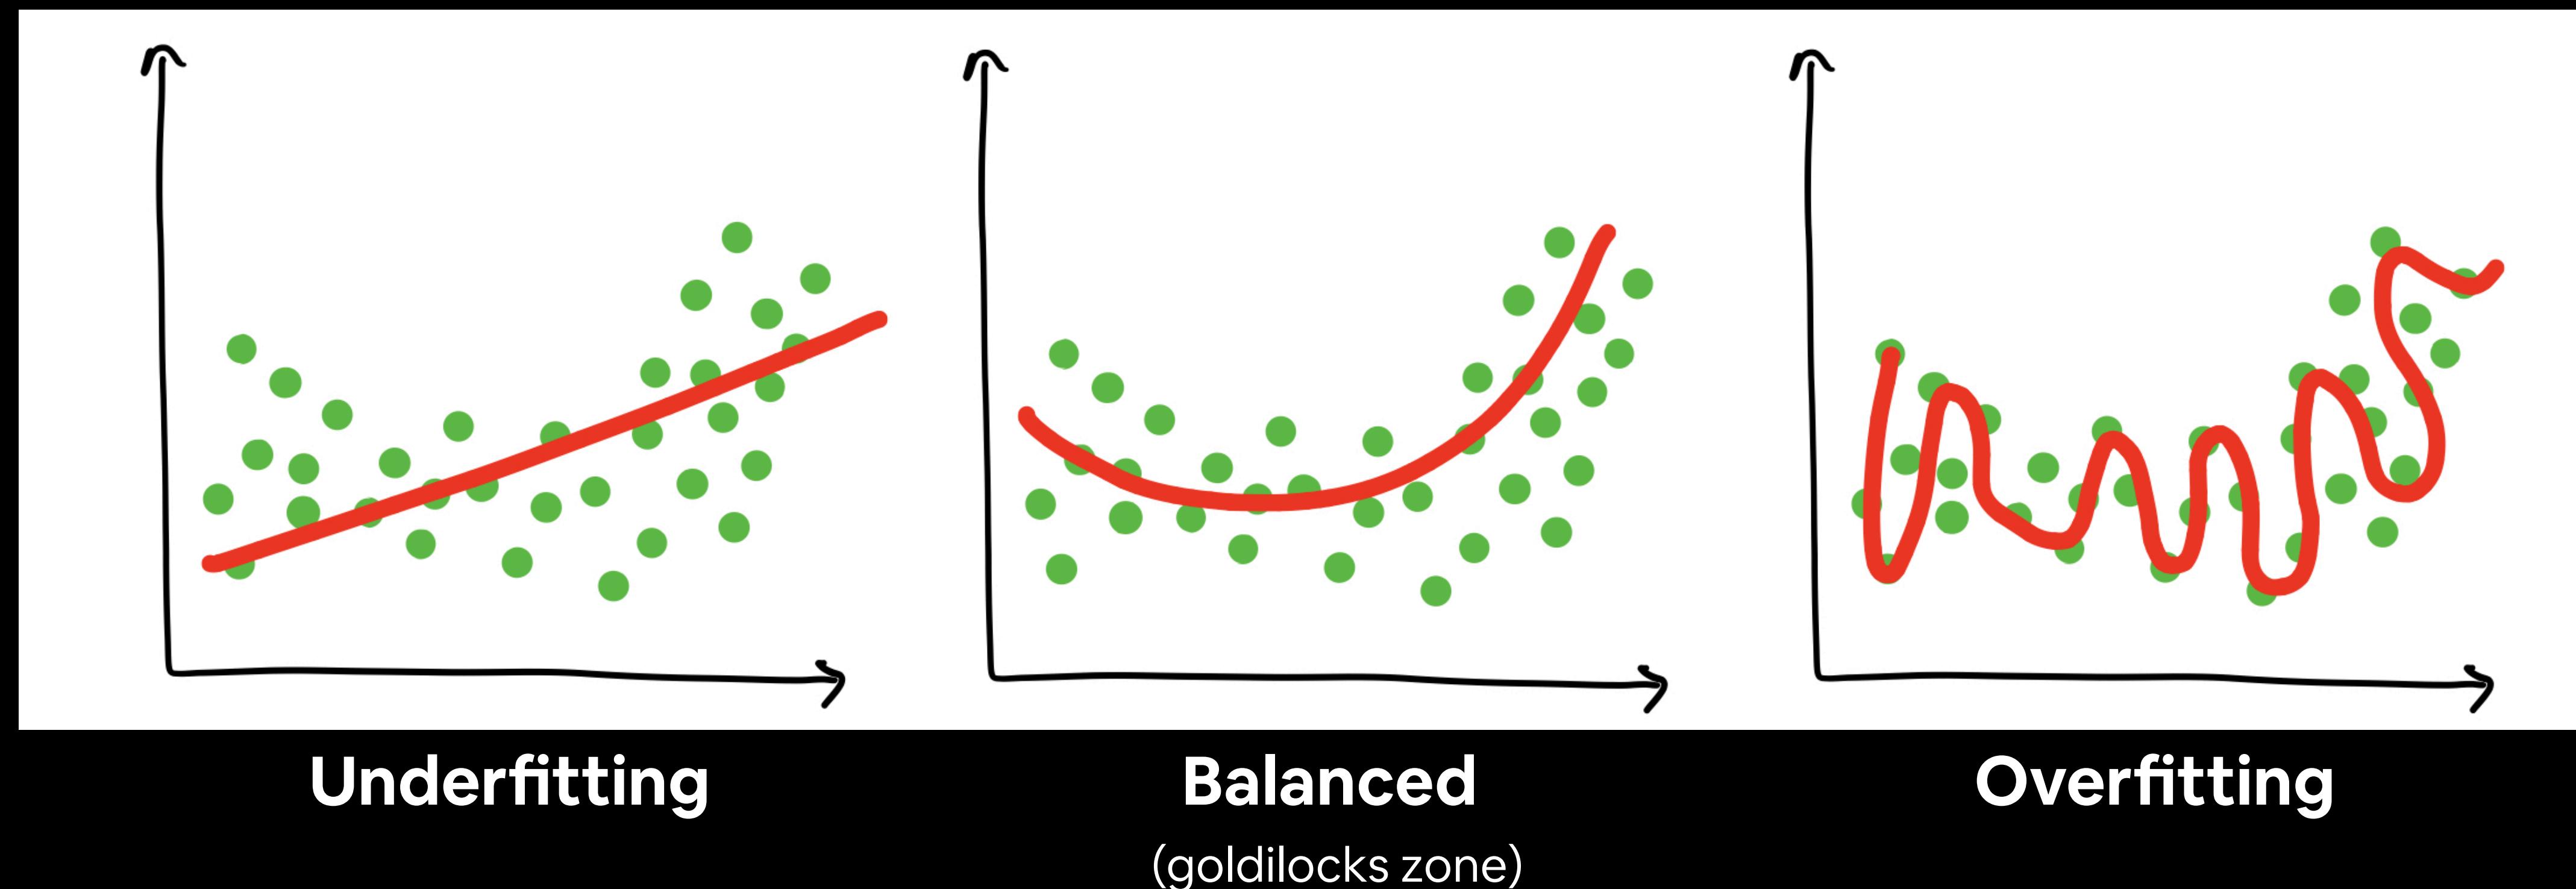
Common ways to improve a deep model:

- Adding layers
- Increase the number of hidden units
- Change the activation functions
- Change the optimization function
- Change the learning rate (because you can alter each of these, they're hyperparameters)
- Fitting on more data
- Fitting for longer

What is overfitting?

Overfitting — when a model over learns patterns in a particular dataset and isn't able to generalise to unseen data.

For example, a student who studies the course materials too hard and then isn't able to perform well on the final exam. Or tries to put their knowledge into practice at the workplace and finds what they learned has nothing to do with the real world.



Improving a model

(from a data perspective)

Method to improve a model (reduce overfitting)

What does it do?

More data

Gives a model more of a chance to learn patterns between samples (e.g. if a model is performing poorly on images of pizza, show it more images of pizza).

Data augmentation (usually for images)

Increase the diversity of your training dataset without collecting more data (e.g. take your photos of pizza and randomly rotate them 30°). Increased diversity forces a model to learn more generalisation patterns.

Better data

Not all data samples are created equally. Removing poor samples from or adding better samples to your dataset can improve your model's performance.

Use transfer learning

Take a model's pre-learned patterns from one problem and tweak them to suit your own problem. For example, take a model trained on pictures of cars to recognise pictures of trucks.

The machine learning explorer's motto

“Visualize, visualize, visualize”



The machine learning practitioner's motto

“Experiment, experiment, experiment”



*(try lots of things and see what
tastes good)*